



MITx 6.86x

Machine Learning with Python-From Linear Models to Deep Learning[Course](#)[Progress](#)[Dates](#)[Discussion](#)[Resources](#)[Course](#) / [Unit 2. Nonlinear Classification, Linear regression,...](#) / [Project 2: Dig](#)[< Previous](#)

2. Linear Regression with Closed Form Solution

[Bookmark this page](#)

Project due Mar 15, 2023 08:59 -03 Completed

After seeing the problem, your classmate Alice immediately argues that we can apply a linear regression model as the labels are numbers from 0-9, very similar to the example we learned from Unit 1. However, being a bit doubtful, you decide to have a try and start simple by using the raw pixel values of each image as features.

Alice wrote a skeleton code `run_linear_regression_on_MNIST` in `main.py`, but she needs your help to complete the code and make the model work.

Closed Form Solution of Linear Regression

5.0/5.0 points (graded)

To solve the linear regression problem, you recall the linear regression has a closed form solution:

$$\theta = (X^T X + \lambda I)^{-1} X^T Y$$

where I is the identity matrix.

Write a function `closed_form` that computes this closed form solution given the feature matrix `X` and the target vector `Y`, and the regularization parameter λ .

Available Functions: You have access to the NumPy python library as `np`; No need to import it.

```
1 def closed_form(X, Y, lambda_factor):
2     """
3     Computes the closed form solution of linear regression with L2 regularization.
4
5     Args:
6         X - (n, d + 1) NumPy array (n datapoints each with d features plus the y-axis intercept)
7         Y - (n, ) NumPy array containing the labels (a number from 0-9) for each data point
8         lambda_factor - the regularization constant (scalar)
9
10    Returns:
11        theta - (d + 1, ) NumPy array containing the weights of linear regression
12        theta[0] - represents the y-axis intercept of the model and therefore X[0] = 1 for all data points
13    """
14    theta = np.linalg.inv(X.T @ X + lambda_factor * np.identity(X.shape[1])) @ X.T @ Y
15    return theta
```

Press ESC then TAB or click outside of the code editor to exit

Unanswered

Submit

You have used 2 of 25 attempts

0.7698

✓

0.7702

✓

Submit

You have used 1 of 20 attempts

What went Wrong?

0.0/1.0 point (graded)

Alice and you find that no matter what factor you try, the test error is large. With some thought, you realize that something is wrong with this approach.



Gradient descent should be used instead of the closed form solution.



The loss function related to the closed-form solution is inadequate for this problem.



Regularization should not be used here.



Submit

You have used 2 of 2 attempts

[< Previous](#)[Next >](#)

Discussion

**edX**[About](#)[Affiliates](#)[edX for Business](#)[Open edX](#)[Careers](#)

[hi everyone, I completed the closed_form method and tried to run main.py and i was getting this error. numpy](#)

? [Memory error](#)

[Hi everyone! I was able to complete the first part of the question, but when I try to run main.py for the 2nd p](#)

💬 [package dependencies issue](#)

[I downloaded the files and ran python main.py. But I'm getting the following message: "qt.qpa.plugin: Could n](#)

💬 [How did we get this formula for the closed solution?](#)

? [TypeError related to transpose](#)

Connect

[Blog](#)

[Contact Us](#)

[Help Center](#)

[Security](#)

[Media Kit](#)



© 2023 edX LLC. All rights reserved.

深圳市恒宇博科技有限公司 [粤ICP备17044299号-2](#)