2. Q-Value Iteration

Homework due May 3, 2023 08:59 -03 Completed

Consider an Markov Decision Process with 6 states $s \in \{0, 1, 2, 3, 4, 5\}$ and 2 actions by the following transition probability functions

For states 1, 2, and 3:

$$T(s, M, s-1) = 1$$

$$T(s, C, s+2) = 0.7$$

$$T(s, C, s) = 0.3$$

For state 0:

$$T(s, M, s) = 1$$

$$T(s, C, s) = 1$$

For states 4 and 5:

$$T(s, M, s-1) = 1$$

$$T(s, C, s) = 1$$

Note that all transition probabilities not defined by the above are equal to 0.

The rewards R are defined by:

$R(s, a, s') = |s' - s|^{\frac{1}{3}} \ \forall s \neq s'$,

and $R(s, a, s) = (s+4)^{\frac{-1}{2}}$, $\forall s \neq 0$.

$R(0, M, 0) = R(0, C, 0) = 0$. Also, the discount factor $\gamma = 0.6$.

We initialize $Q_0(s, a) = 0 \forall s \in \{0, 1, 2, 3, 4, 5\}$ and $\forall a \in \{C, M\}$

2

6.0/6.0 points (graded)

Input the Q-values **correct to 3 decimal places** after one Q-value iteration

0	✔
0	✔
1	✔
1.016	✔
1	✔
1.004	✔
1	✔
0.995	✔
1	✔
0.354	✔
1	✔
0.333	✔

Submit

You have used 1 of 4 attempts

1

Submit You have used 1 of 2 attempts

4

5/5 points (graded)

What are the optimal policies we get from ?

- C
- M

- C
- M

Previous Next

Show all posts

? Reward clarification

For the case where both s != s' and s != 0 hold true, which reward function should we use? The first, the sec

Connect

Blog
Contact Us
Help Center
Security
Media Kit