



MITx 6.86x

Machine Learning with Python-From Linear Models to Deep Learning[Course](#)[Progress](#)[Dates](#)[Discussion](#)[Resources](#)[Course](#) / [Unit 5. Reinforcement Learning \(2 weeks\)](#) / [Project 5: Text-Based Game](#)[< Previous](#)

6. Q-learning with linear function approximation

[Bookmark this page](#)

Project due May 10, 2023 08:59 -03 Completed

Since the state displayed to the agent is described in text, we have to choose a mechanism to convert textual descriptions into vector representations. A naive way is to create one unique index for each state-action pair we have done in previous part. However, such approach becomes infeasible when the state-action space is huge. To tackle this challenge, we can design some representation generator that does not depend on the textual state space. In particular, a representation generator $\psi_R(\cdot)$ reads raw text description and converts it to a vector representation $\mathbf{v}_s = \psi_R(s)$. One approach is to use a bag-of-words model derived from the text description.

In large games, it is often impractical to maintain the Q-value for all possible state-action pairs. One solution to this problem is to approximate $Q(s, c)$ using a parametrized function $Q(s, c; \theta)$.

In this section we consider a linear parametric architecture:

$$Q(s, c; \theta) = \phi(s, c)^T \theta = \sum_{i=1}^d \phi_i(s, c) \theta_i,$$

where $\phi(s, c)$ is a fixed feature vector in \mathbb{R}^d for state-action pair (s, c) with i -th component $\phi_i(s, c)$, and $\theta \in \mathbb{R}^d$ is a parameter vector that is shared across state-action pairs. The key challenge is to design the feature vectors $\phi(s, c)$. Note that given a textual state s , we first translate it to a vector representation using $\psi_R(s)$. So the question here is how to design a mapping function convert $(\psi_R(s), c)$ to a feature vector in \mathbb{R}^d . Assume that the size of action space is d_C , and the dimension of state representation is d_R .

Feature engineering

1/1 point (graded)

Exercise: Consider the following feature engineering. Define a function $\psi_C : \mathcal{C} \rightarrow \mathbb{R}^{d_C}$ where the j -th component $\psi_{C,j}(c)$ is given as follows:

$$\psi_{C,j}(c) = \begin{cases} 1 & \text{if } j = c \\ 0 & \text{else} \end{cases}$$

The feature vector is defined as

Alternatively, consider the following feature map:

, where

, and for

,

$$\begin{bmatrix} 1 & 0 & 1 \\ \vdots & & \\ 1 & 0 & 1 \end{bmatrix}$$

You will implement this feature map in the next tab.

Computing theta update rule

1.0/1 point (graded)

The Q-learning approximation algorithm starts with an initial parameter estimate of θ . A

Q-learning, upon observing a data tuple (s, c, r) , the target value $Q^*(s, c)$ for the C

defined as the sampled version of the Bellman operator,

Then the parameter θ is simply updated by taking a gradient step with respect to the s

< Previous

Next >

The negative gradient can be computed as follows.

(Enter your answer in terms of y , $Q(s, c, \theta)$, and $\phi(s, c)$.)



edX

[About](#)

[Affiliates](#)

in the alternative feature map there seems to be a problem displaying a part of the text: Could not format H

🗨 Invalid Input: '\d\' not permitted in answer as a variable

Hi all, I get the messageInvalid Input: '\d\' not permitted in answer as a variable even when I have no 'd' in my

? What is the shape of the feature vector/matrix?

? Why is there no "Show answer" here?

I got the answer wrong for the first question here and while I think I understand now why it is wrong, a detail

? Specifics of the feature array_phi

Is the structure of the feature array such that for each state description, only a range of the rows have nonze

☑ $y = R(s,c) + \gamma * Q(s', c')$?

Connect

[Blog](#)

[Contact Us](#)

[Help Center](#)

[Security](#)

[Media Kit](#)



© 2023 edX LLC. All rights reserved.

深圳市恒宇博科技有限公司 [粤ICP备17044299号-2](#)